

✓ PART-1: Dense Network

```
%load_ext tensorboard
import tensorflow as tf
import datetime

# Clear any logs from previous runs
!rm -rf ./logs/

%reload_ext tensorboard

import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from torchvision import datasets, transforms
import numpy as np
from google.colab import drive
from torch.utils.tensorboard import SummaryWriter # Corrected import statement

# Load the data from the file
data = np.load('emnist_letters.npz')

# Access the arrays containing images and labels
train_images = data['train_images']
train_labels = data['train_labels']
val_images = data['validate_images']
val_labels = data['validate_labels']
test_images = data['test_images']
test_labels = data['test_labels']

# Create datasets
train_dataset = TensorDataset(torch.tensor(train_images), torch.tensor(train_labels))
val_dataset = TensorDataset(torch.tensor(val_images), torch.tensor(val_labels))
test_dataset = TensorDataset(torch.tensor(test_images), torch.tensor(test_labels))

# Create data loaders
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=64)
```

```
test_loader = DataLoader(test_dataset, batch_size=64)
```

```
device = (
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.backends.mps.is_available()
    else "cpu"
)
print(f"Using {device} device")
```


 Using cuda device

```
# Model Definition
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        # self.flatten = nn.Flatten()
        self.fc1 = nn.Linear(28 * 28, 300, dtype=torch.float32) # One neuron per fe
        self.fc2 = nn.Linear(300,150, dtype=torch.float32) # arbitrarily selecting
        self.fc3 = nn.Linear(150, 50, dtype=torch.float32)
        self.fc4 = nn.Linear(50, 27, dtype=torch.float32) # 27 Classes present in d

    def forward(self, x):
        # x = self.flatten(x)
        # Cast input data to torch.float32
        x = x.to(torch.float32)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = self.fc4(x)
        return x
```

```
model = NeuralNetwork()
print(model)
```

```
# Inspect a sample of the data
for data, target in train_loader:
    print("Target shape:", target.shape)
    print("Target dataset:", target)
    break
```

 NeuralNetwork(
 (fc1): Linear(in_features=784, out_features=300, bias=True)
 (fc2): Linear(in_features=300, out_features=150, bias=True)

```

(fc3): Linear(in_features=150, out_features=50, bias=True)
(fc4): Linear(in_features=50, out_features=27, bias=True)
)
Target shape: torch.Size([64, 27])
Target dataset: tensor([[0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 1., 0.],
                        ...,
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 1., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 1., 0., 0.]])

```

```
# Training
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()
```

```
# Create a directory for TensorBoard logs
```

```
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
os.makedirs(log_dir, exist_ok=True)
```

```
# Create a SummaryWriter for TensorBoard
```

```
writer = SummaryWriter(log_dir=log_dir)
```

```
def train(model, train_loader, optimizer, criterion, epochs):
```

```
    for epoch in range(epochs):
```

```
        model.train()
```

```
        epoch_loss = 0.0 # Track epoch loss
```

```
        for batch_idx, (data, target) in enumerate(train_loader):
```

```
            optimizer.zero_grad()
```

```
            output = model(data)
```

```
            # Convert one-hot encoded target to class labels (1D tensor)
```

```
            target_labels = torch.nonzero(target, as_tuple=True)[1]
```

```
            # Calculate loss using class labels
```

```
            loss = criterion(output, target_labels)
```

```
            loss.backward()
```

```
            optimizer.step()
```

```
        epoch_loss += loss.item() # Accumulate batch loss
```

```
# Print epoch results
```

```
print('Epoch {} - Loss: {:.6f}'.format(epoch, epoch_loss / len(train_loader)))
```

```
# Log the training loss to TensorBoard
```

```
writer.add_scalar('Loss/train', epoch_loss / len(train_loader), epoch)
```

```
def test(model, test_loader):
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            output = model(data)
            # Convert one-hot encoded target to class labels (1D tensor)
            target_labels = torch.nonzero(target, as_tuple=True)[1]
            test_loss += criterion(output, target_labels).item()
            pred = output.argmax(dim=1, keepdim=True)
            correct += pred.eq(target_labels.view_as(pred)).sum().item()

    test_loss /= len(test_loader.dataset)
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%) \n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))

train(model, train_loader, optimizer, criterion, epochs=10)
# Close the SummaryWriter
writer.close()
test(model, test_loader)
```

```
⇒ Epoch 0 - Loss: 0.835394
Epoch 1 - Loss: 0.388811
Epoch 2 - Loss: 0.306091
Epoch 3 - Loss: 0.261132
Epoch 4 - Loss: 0.230223
Epoch 5 - Loss: 0.204188
Epoch 6 - Loss: 0.187152
Epoch 7 - Loss: 0.171736
Epoch 8 - Loss: 0.158153
Epoch 9 - Loss: 0.146823
```

```
Test set: Average loss: 0.0050, Accuracy: 18932/20800 (91%)
```

✓ Graphing the data from the Dense network

First, let's generate confusion matrix of the dense network's classifications

We'll start by calculating the number of true and false positives

```
import numpy as np

# Generate predictions for the test set
def generate_predictions(model, test_loader):
    model.eval()
```

```

all_predictions = []
all_targets = []
with torch.no_grad():
    for data, target in test_loader:
        output = model(data)
        all_predictions.extend(output.argmax(dim=1).cpu().numpy())
        all_targets.extend(target.cpu().numpy())
return np.array(all_predictions), np.array(all_targets)
def compute_tp_fp(predictions, targets, class_label):
    # Convert one-hot encoded targets to class labels
    target_labels = np.argmax(targets, axis=1)

    # Compute True Positives (TP) and False Positives (FP) for the specified class l
    tp = np.sum((predictions == class_label) & (target_labels == class_label))
    fp = np.sum((predictions == class_label) & (target_labels != class_label))

    return tp, fp

# Generate predictions for the test set
test_predictions, test_targets = generate_predictions(model, test_loader)

# Compute TP and FP for each class
for class_label in range(27):
    tp, fp = compute_tp_fp(test_predictions, test_targets, class_label)
    print(f"Class {class_label}: TP={tp}, FP={fp}")

```



```

Class 0: TP=0, FP=0
Class 1: TP=699, FP=77
Class 2: TP=765, FP=87
Class 3: TP=749, FP=53
Class 4: TP=737, FP=73
Class 5: TP=758, FP=73
Class 6: TP=744, FP=55
Class 7: TP=615, FP=123
Class 8: TP=721, FP=64
Class 9: TP=511, FP=138
Class 10: TP=751, FP=69
Class 11: TP=730, FP=54
Class 12: TP=645, FP=279
Class 13: TP=769, FP=37
Class 14: TP=763, FP=107
Class 15: TP=753, FP=47
Class 16: TP=760, FP=36
Class 17: TP=666, FP=179
Class 18: TP=729, FP=61
Class 19: TP=758, FP=29
Class 20: TP=765, FP=78
Class 21: TP=730, FP=57
Class 22: TP=704, FP=36
Class 23: TP=752, FP=17
Class 24: TP=752, FP=36
Class 25: TP=744, FP=69
Class 26: TP=770, FP=26

```

✓ Confusion Matrix

```
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.metrics import confusion_matrix


# Generate predictions for the test set
test_predictions, test_targets = generate_predictions(model, test_loader)

# Convert one-hot encoded targets to class labels
test_targets_single = np.argmax(test_targets, axis=1)

# Compute confusion matrix
cm = confusion_matrix(test_targets_single, test_predictions)

# Generate labels for the letters using Unicode
letter_labels = [chr(ord('A') + i) for i in range(26)]

# Plot confusion matrix as heatmap with x-axis on top and "plasma" colormap
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="plasma", xticklabels=letter_labels, ytick
plt.xlabel("Predicted Letter")
plt.ylabel("True Letter")
plt.title("Confusion Matrix")
plt.show()
```



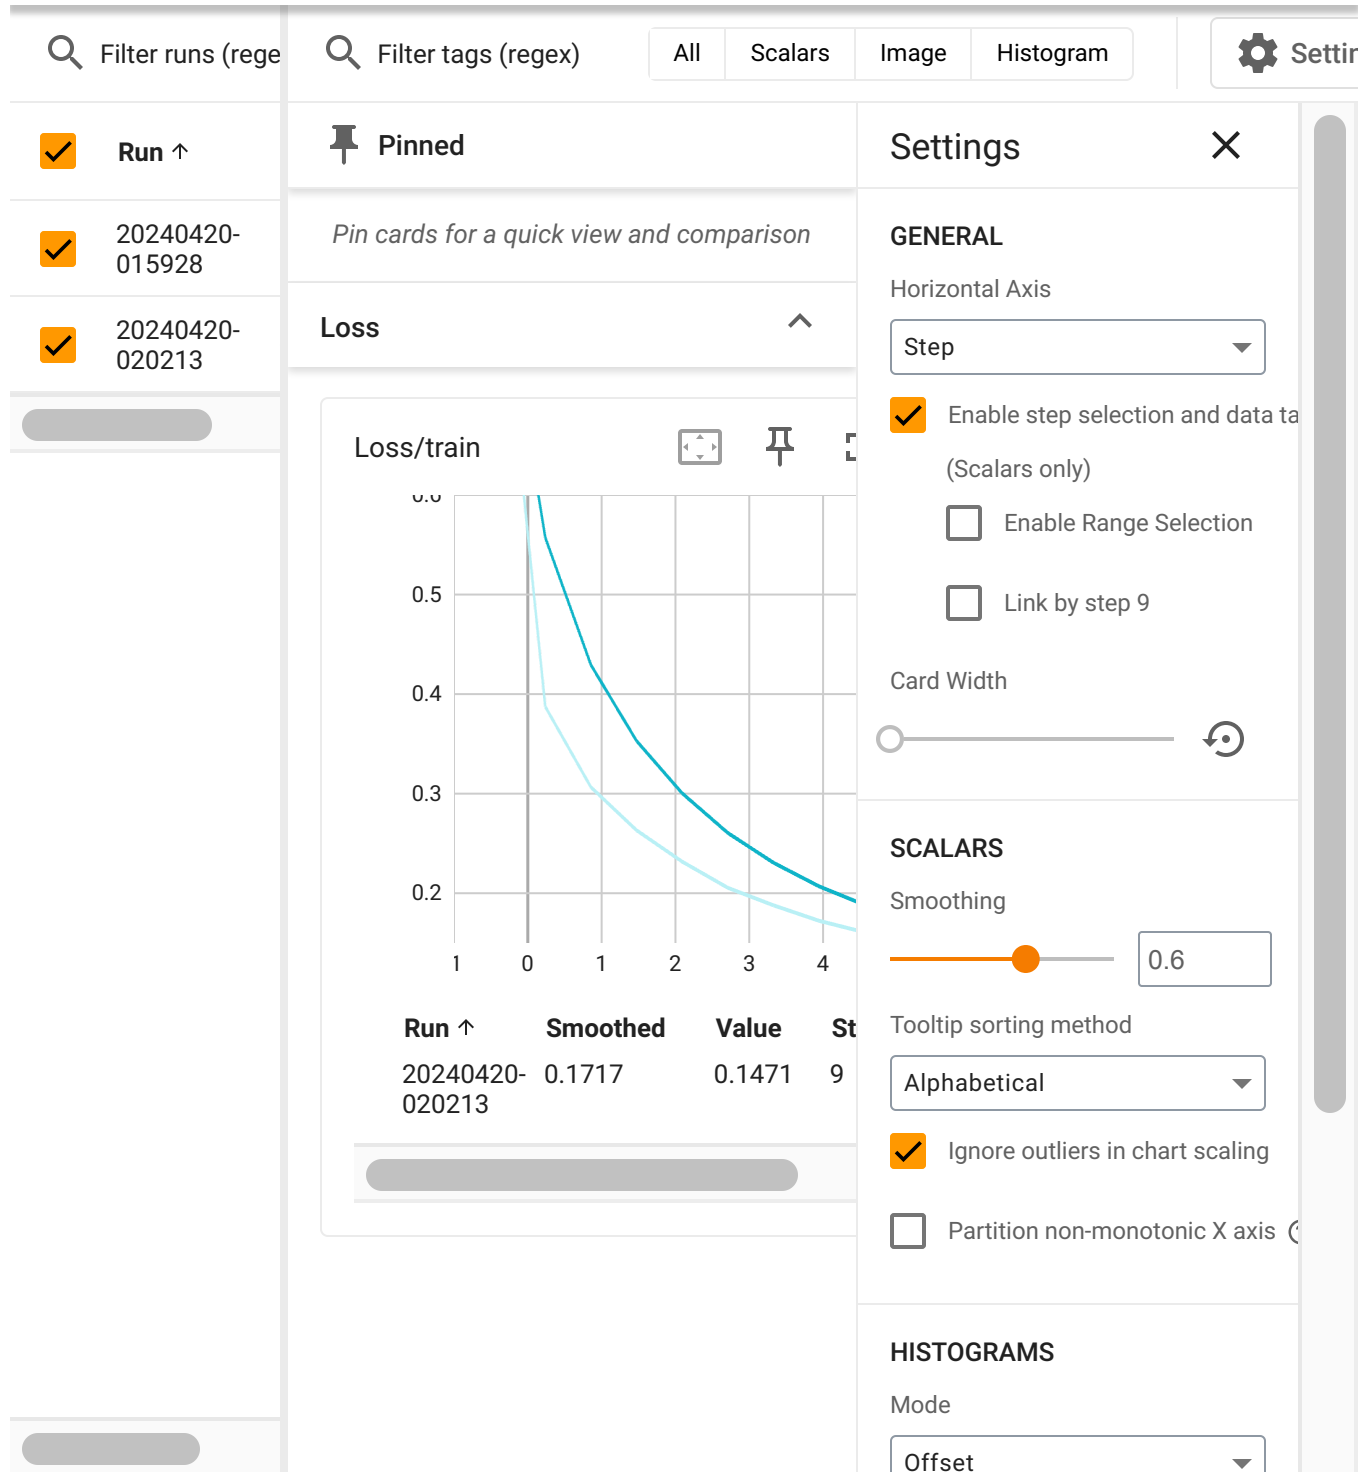
```
%tensorboard --logdir logs/fit
```



TensorBoard

TIME SERIES

SCALARS INACTIVE



Performance Comparison between Dense Network and OPIUM based Classifier

The Dense Network used 3 hidden layers but the OPIUM based classifier used 10,000 hidden layers. Even with the huge increase in the hidden layers the accuracy of the OPIUM based classifier on the letters dataset remained at $(85.15\% \pm 0.12\%)$, while the dense network had a much better accuracy of 91%. In comparison with the OPIUM based classifier the dense network is more compact and runs more efficiently.

✓ PART-2: Convolutional Network

```
import tensorflow as tf
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, TensorBoard
import datetime
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Activation, BatchNormalization, MaxPooli
from tensorflow.keras.models import load_model
```

```
# Clear any logs from previous runs
!rm -rf ./logs/
# Load the data from the file
data = np.load('emnist_letters.npz')
```

```
# Access the arrays containing images and labels
train_images = data['train_images']
train_labels = data['train_labels']
validate_images = data['validate_images']
validate_labels = data['validate_labels']
test_images = data['test_images']
test_labels = data['test_labels']
```

```
# Define the log directory for TensorBoard
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
# Load the TensorBoard notebook extension
%load_ext tensorboard
```

```
# Define TensorBoard callback
tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1)
```



The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

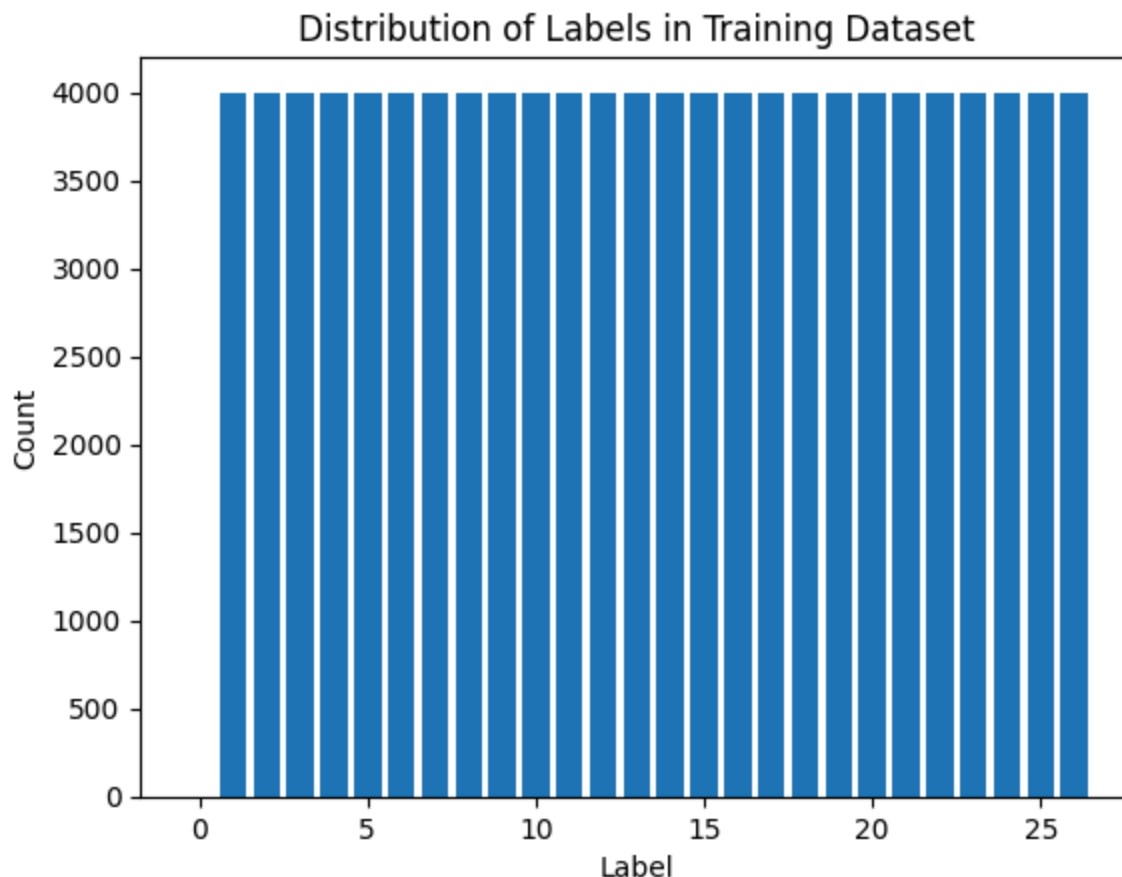
```
num_train_images = train_images.shape[0]
print("Number of images in the training dataset:", num_train_images)
```

⇒ Number of images in the training dataset: 104000

```
# Count occurrences of each label
label_counts = np.sum(train_labels, axis=0)
```

```
# Plot the distribution of labels
plt.bar(range(len(label_counts)), label_counts)
plt.xlabel('Label')
plt.ylabel('Count')
plt.title('Distribution of Labels in Training Dataset')
plt.show()
```

⇒



```
print(train_images.shape)
print(validate_images.shape)
```

⇒ (104000, 784)
(20800, 784)

```
!pip install tensorflow
```

⇒ Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-

```

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: flatbuffers>=23.5.26 in /usr/local/lib/python3.10
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dis
Requirement already satisfied: ml-dtypes~0.2.0 in /usr/local/lib/python3.10/dis
Requirement already satisfied: numpy<2.0.0,>=1.23.5 in /usr/local/lib/python3.10
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/di
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.2
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dis
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python
Requirement already satisfied: wrapt<1.15,>=1.11.0 in /usr/local/lib/python3.10/
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/loca
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/
Requirement already satisfied: tensorboard<2.16,>=2.15 in /usr/local/lib/python3
Requirement already satisfied: tensorflow-estimator<2.16,>=2.15.0 in /usr/local/
Requirement already satisfied: keras<2.16,>=2.15.0 in /usr/local/lib/python3.10/
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.1
Requirement already satisfied: google-auth-oauthlib<2,>=0.5 in /usr/local/lib/py
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/loc
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.1
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/d
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/d
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: pyasn1<0.7.0,>=0.4.6 in /usr/local/lib/python3.10
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.10/dist

```

```

# Define the strategy
strategy = tf.distribute.MirroredStrategy()
with strategy.scope():
    model = Sequential()

# Reshape the input images to their original 28x28 shape (assuming original shap
model.add(tf.keras.layers.Reshape((28, 28, 1), input_shape=(784,)))

# Feature Learning Layers
model.add(Conv2D(32,                # Number of filters/Kernels
                 (3,3),              # Size of kernels (3x3 matrix)
                 strides = 1,         # Step size for sliding the kernel across

```

```

        padding = 'same'      # 'Same' ensures that the output feature ma
    ))
    model.add(Activation('relu'))# Activation function
    model.add(BatchNormalization())
    model.add(MaxPooling2D(pool_size = (2,2), padding = 'same'))
    model.add(Dropout(0.2))

    model.add(Conv2D(64, (5,5), padding = 'same'))
    model.add(Activation('relu'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D(pool_size = (2,2), padding = 'same'))
    model.add(Dropout(0.2))

    model.add(Conv2D(128, (3,3), padding = 'same'))
    model.add(Activation('relu'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D(pool_size = (2,2), padding = 'same'))
    model.add(Dropout(0.3))

    # Flattening tensors
    model.add(Flatten())

    # Fully-Connected Layers
    model.add(Dense(2048))
    model.add(Activation('relu'))
    model.add(Dropout(0.5))

    # Output Layer
    model.add(Dense(27, activation = 'softmax')) # Classification layer

model.compile(optimizer = tf.keras.optimizers.RMSprop(0.0001), # 1e-4
              loss = 'categorical_crossentropy', # Ideal for multiclass tasks
              metrics = ['accuracy']) # Evaluation metric

# Defining an Early Stopping and Model Checkpoints
early_stopping = EarlyStopping(monitor = 'val_accuracy',
                               patience = 5, mode = 'max',
                               restore_best_weights = True)

checkpoint = ModelCheckpoint('best_model.h5',
                             monitor = 'val_accuracy',
                             save_best_only = True)

# Define the number of epochs
num_epochs = 50

# Fit the model to the training data
history = model.fit(train_images, train_labels,
                    epochs=num_epochs,
                    validation_data=(validate_images, validate_labels),

```

```
callbacks=[early_stopping, checkpoint, tensorboard_callback])
```

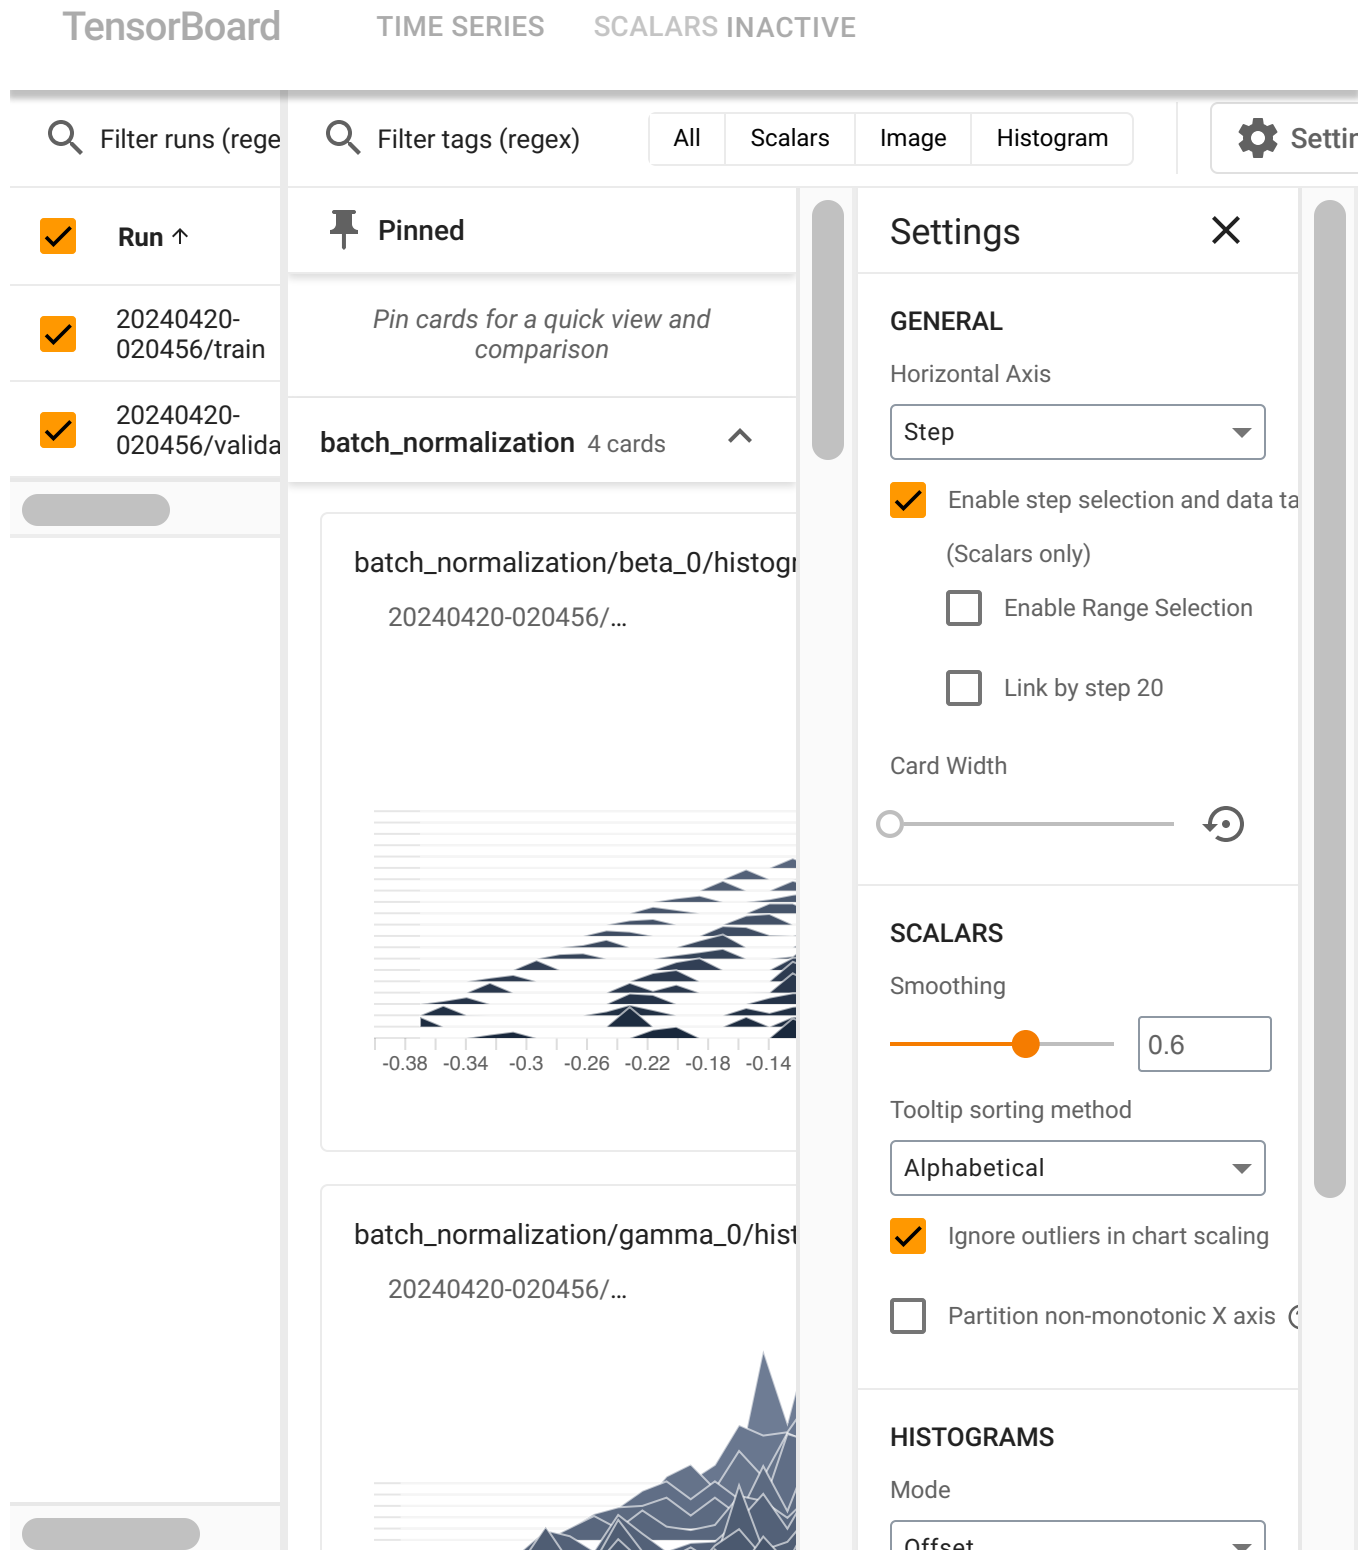


```
Epoch 1/50
3248/3250 [=====>.] - ETA: 0s - loss: 0.7824 - accuracy:
  saving_api.save_model(
3250/3250 [=====] - 37s 8ms/step - loss: 0.7821 - accur
Epoch 2/50
3250/3250 [=====] - 27s 8ms/step - loss: 0.3694 - accur
Epoch 3/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.3162 - accur
Epoch 4/50
3250/3250 [=====] - 28s 9ms/step - loss: 0.2871 - accur
Epoch 5/50
3250/3250 [=====] - 26s 8ms/step - loss: 0.2725 - accur
Epoch 6/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2670 - accur
Epoch 7/50
3250/3250 [=====] - 26s 8ms/step - loss: 0.2585 - accur
Epoch 8/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2532 - accur
Epoch 9/50
3250/3250 [=====] - 27s 8ms/step - loss: 0.2506 - accur
Epoch 10/50
3250/3250 [=====] - 26s 8ms/step - loss: 0.2447 - accur
Epoch 11/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2402 - accur
Epoch 12/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2404 - accur
Epoch 13/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2337 - accur
Epoch 14/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2322 - accur
Epoch 15/50
3250/3250 [=====] - 27s 8ms/step - loss: 0.2253 - accur
Epoch 16/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2241 - accur
Epoch 17/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2201 - accur
Epoch 18/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2183 - accur
Epoch 19/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2146 - accur
Epoch 20/50
3250/3250 [=====] - 26s 8ms/step - loss: 0.2164 - accur
Epoch 21/50
3250/3250 [=====] - 25s 8ms/step - loss: 0.2114 - accur
```



```
%tensorboard --logdir logs/fit
```

➔ Reusing TensorBoard on port 6006 (pid 3879), started 0:10:45 ago. (Use '!kill 3879' to kill it.)



```
# Load the best model
best_model = load_model('best_model.h5')

# Evaluate the best model on test data
test_loss, test_accuracy = best_model.evaluate(test_images, test_labels)
```

```
print('Test Loss:', test_loss)
print('Test Accuracy:', test_accuracy)
```

```
650/650 [=====] - 2s 3ms/step - loss: 0.1834 - accuracy
Test Loss: 0.18337306380271912
Test Accuracy: 0.9402884840965271
```

Generate True and False Positives for Convolutional Network

```
import numpy as np
from tensorflow.keras.models import load_model

# Load the best model
best_model = load_model('best_model.h5')

# Generate predictions for the test set
def generate_predictions(model, test_images):
    predictions = model.predict(test_images)
    return np.argmax(predictions, axis=1)

# Load the test data
test_images = data['test_images']
test_labels = data['test_labels']

# Flatten the test labels
test_labels_flat = np.argmax(test_labels, axis=1)

# Compute TP and FP for each class
def compute_tp_fp(predictions, targets, class_label):
    # Compute True Positives (TP) and False Positives (FP) for the specified class
    tp = np.sum((predictions == class_label) & (targets == class_label))
    fp = np.sum((predictions == class_label) & (targets != class_label))
    return tp, fp

# Generate predictions for the test set
test_predictions = generate_predictions(best_model, test_images)

# Compute TP and FP for each class
for class_label in range(27):
    tp, fp = compute_tp_fp(test_predictions, test_labels_flat, class_label)
    print(f"Class {class_label}: TP={tp}, FP={fp}")
```

```
650/650 [=====] - 1s 2ms/step
Class 0: TP=0, FP=0
Class 1: TP=776, FP=74
Class 2: TP=774, FP=21
Class 3: TP=782, FP=25
Class 4: TP=746, FP=37
```

```
Class 5: TP=773, FP=16
Class 6: TP=769, FP=8
Class 7: TP=674, FP=101
Class 8: TP=757, FP=31
Class 9: TP=637, FP=249
Class 10: TP=737, FP=12
Class 11: TP=769, FP=10
Class 12: TP=580, FP=177
Class 13: TP=794, FP=16
Class 14: TP=774, FP=44
Class 15: TP=792, FP=69
Class 16: TP=787, FP=30
Class 17: TP=673, FP=82
Class 18: TP=765, FP=17
Class 19: TP=789, FP=24
Class 20: TP=780, FP=24
Class 21: TP=762, FP=52
Class 22: TP=750, FP=57
Class 23: TP=789, FP=14
Class 24: TP=779, FP=24
Class 25: TP=756, FP=21
Class 26: TP=794, FP=7
```


✓ Confusion Matrix

```
import matplotlib.pyplot as plt
import seaborn as sns

# Define class labels (assuming class labels are represented as integers from 0 to 26)
class_labels = range(26)
letter_labels = [chr(ord('A') + i) for i in class_labels]

# Compute confusion matrix
cm = confusion_matrix(test_labels_flat, test_predictions)

# Plot confusion matrix as heatmap with "viridis" colormap
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="viridis", xticklabels=letter_labels, yticklabels=letter_labels)
plt.xlabel("Predicted Letter")
plt.ylabel("True Letter")
plt.title("Confusion Matrix")
plt.show()
```

Performance Comparison between Dense Network and CNN

The dense network performed at 91% accuracy while the convolution network performed a bit better at 94%. The models both struggled in the same areas and both had most of their misidentifications in the same place. Similar letters were frequently misidentified as each other- for example both of the networks had the most issues misidentifying the letter I as the letter L, and vice versa. The second highest misidentifications were Q and G. Overall though, the Convolution

network was more consistent in identifying letters, with far more pairs of letters at 0 total misidentifications.

✓ Part-3: GAN

```
import numpy as np
from torch.utils.data import Dataset, DataLoader
import argparse
import os
import random
import torch
import torch.nn as nn
import torch.nn.parallel
import torch.optim as optim
import torch.utils.data
import torchvision.datasets as dset
import torchvision.transforms as transforms
import torchvision.utils as vutils
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.animation as animation
from IPython.display import HTML
from torch.utils.tensorboard import SummaryWriter
writer = SummaryWriter()

# Load the data from the file

data = np.load('emnist_letters.npz')

# Access the arrays containing images and labels
train_images = data['train_images']
train_labels = data['train_labels']
validate_images = data['validate_images']
validate_labels = data['validate_labels']
test_images = data['test_images']
test_labels = data['test_labels']

# Concatenate images and labels arrays
all_images = np.concatenate([train_images, validate_images, test_images], axis=0)
all_labels = np.concatenate([train_labels, validate_labels, test_labels], axis=0)

# Number of workers for dataloader
workers = 2

# Batch size during training
```

```
batch_size = 128

# Spatial size of training images. All images will be resized to this
# size using a transformer.
image_size = 64

# Number of channels in the training images. For color images this is 3
nc = 1

# Size of z latent vector (i.e. size of generator input)
nz = 100

# Size of feature maps in generator
ngf = 28

# Size of feature maps in discriminator
ndf = 28

# Number of training epochs
num_epochs = 50

# Learning rate for optimizers
lr = 0.0002

# Beta1 hyperparameter for Adam optimizers
beta1 = 0.5

# Number of GPUs available. Use 0 for CPU mode.
ngpu = 1


class CustomDataset(Dataset):
    def __init__(self, images, labels, transform=None):
        self.images = images
        self.labels = labels
        self.transform = transform

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        image = self.images[idx].reshape(28, 28) # Reshape flattened image to 2D
        label = self.labels[idx]

        if self.transform:
            image = self.transform(image)

        return image, label

# Transform for image preprocessing
```

```
transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize(image_size),
    transforms.CenterCrop(image_size),
    transforms.ToTensor()
])

# Create custom dataset instances
train_dataset = CustomDataset(all_images, all_labels, transform=transform)

# Create dataloaders
train_dataloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=4)

# Decide which device we want to run on
device = torch.device("cuda:0" if (torch.cuda.is_available() and torch.cuda.device_count() > 0) else "cpu")
real_batch = next(iter(train_dataloader))
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Training Images")
plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=2, normalize=True), (0, 1, 2)))
plt.show()
```

```

/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork()
  self.pid = os.fork()

```

Training Images



```

# custom weights initialization called on ``netG`` and ``netD``
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)

```

```

# Generator Code

```

```

class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d( nz, ngf * 8, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. ``(ngf*8) x 4 x 4``
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
            # state size. ``(ngf*4) x 8 x 8``
            nn.ConvTranspose2d( ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # state size. ``(ngf*2) x 16 x 16``
            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            # state size. ``(ngf) x 32 x 32``
            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
            nn.Tanh()
            # state size. ``(nc) x 64 x 64``
        )

    def forward(self, input):
        return self.main(input)

```

```

# Create the generator
netG = Generator(ngpu).to(device)

```

```

# Handle multi-GPU if desired
if (device.type == 'cuda') and (ngpu > 1):
    netG = nn.DataParallel(netG, list(range(ngpu)))

```

```

# Apply the ``weights_init`` function to randomly initialize all weights
# to ``mean=0``, ``stdev=0.02``.
netG.apply(weights_init)

```

```

# Print the model
print(netG)

```



```

Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 224, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(224, eps=1e-05, momentum=0.1, affine=True, track_running_st
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(224, 112, kernel_size=(4, 4), stride=(2, 2), padding=(1
    (4): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True, track_running_st

```

```

(5): ReLU(inplace=True)
(6): ConvTranspose2d(112, 56, kernel_size=(4, 4), stride=(2, 2), padding=(1,
(7): BatchNorm2d(56, eps=1e-05, momentum=0.1, affine=True, track_running_sta
(8): ReLU(inplace=True)
(9): ConvTranspose2d(56, 28, kernel_size=(4, 4), stride=(2, 2), padding=(1,
(10): BatchNorm2d(28, eps=1e-05, momentum=0.1, affine=True, track_running_st
(11): ReLU(inplace=True)
(12): ConvTranspose2d(28, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1,
(13): Tanh()
)
)

```

```

class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is ``(nc) x 64 x 64``
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. ``(ndf) x 32 x 32``
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. ``(ndf*2) x 16 x 16``
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. ``(ndf*4) x 8 x 8``
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 8),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. ``(ndf*8) x 4 x 4``
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.main(input)

```

```

# Create the Discriminator
netD = Discriminator(ngpu).to(device)

# Handle multi-GPU if desired
if (device.type == 'cuda') and (ngpu > 1):
    netD = nn.DataParallel(netD, list(range(ngpu)))

# Apply the ``weights_init`` function to randomly initialize all weights
# like this: ``to mean=0, stdev=0.2``.

```

```
netD.apply(weights_init)
```

```
# Print the model
print(netD)
```

```
↗ Discriminator(
  (main): Sequential(
    (0): Conv2d(1, 28, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=F
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(28, 56, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=
    (3): BatchNorm2d(56, eps=1e-05, momentum=0.1, affine=True, track_running_sta
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(56, 112, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias
    (6): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True, track_running_st
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(112, 224, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bia
    (9): BatchNorm2d(224, eps=1e-05, momentum=0.1, affine=True, track_running_st
    (10): LeakyReLU(negative_slope=0.2, inplace=True)
    (11): Conv2d(224, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (12): Sigmoid()
  )
)
```

```
# Initialize the ``BCELoss`` function
criterion = nn.BCELoss()
```

```
# Create batch of latent vectors that we will use to visualize
# the progression of the generator
fixed_noise = torch.randn(64, nz, 1, 1, device=device)
```

```
# Establish convention for real and fake labels during training
real_label = 1.
fake_label = 0.
```

```
# Setup Adam optimizers for both G and D
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))
```

```
# Training Loop
```

```
# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0
```

```
%load_ext tensorboard
```

```
print("Starting Training Loop...")
```



```

# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(train_dataloader, 0):

        #####
        # (1) Update D network: maximize  $\log(D(x)) + \log(1 - D(G(z)))$ 
        #####
        ## Train with all-real batch
        netD.zero_grad()
        # Format batch
        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        label = torch.full((b_size,), real_label, dtype=torch.float, device=device)
        # Forward pass real batch through D
        output = netD(real_cpu).view(-1)
        # Calculate loss on all-real batch
        errD_real = criterion(output, label)
        # Calculate gradients for D in backward pass
        errD_real.backward()
        D_x = output.mean().item()

        ## Train with all-fake batch
        # Generate batch of latent vectors
        noise = torch.randn(b_size, nz, 1, 1, device=device)
        # Generate fake image batch with G
        fake = netG(noise)
        label.fill_(fake_label)
        # Classify all fake batch with D
        output = netD(fake.detach()).view(-1)
        # Calculate D's loss on the all-fake batch
        errD_fake = criterion(output, label)
        # Calculate the gradients for this batch, accumulated (summed) with previous
        errD_fake.backward()
        D_G_z1 = output.mean().item()
        # Compute error of D as sum over the fake and the real batches
        errD = errD_real + errD_fake
        # Update D
        optimizerD.step()

        #####
        # (2) Update G network: maximize  $\log(D(G(z)))$ 
        #####
        netG.zero_grad()
        label.fill_(real_label) # fake labels are real for generator cost
        # Since we just updated D, perform another forward pass of all-fake batch th
        output = netD(fake).view(-1)
        # Calculate G's loss based on this output
        errG = criterion(output, label)
        # Calculate gradients for G
        errG.backward()

```

```

D_G_z2 = output.mean().item()
# Update G
optimizerG.step()

# Output training stats
if i % 50 == 0:
    print('[%d/%d] [%d/%d] \t Loss_D: %.4f \t Loss_G: %.4f \t D(x): %.4f \t D(G(z)):'
          % (epoch, num_epochs, i, len(train_dataloader),
             errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(train_dataloader) - 1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
    img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

# Log scalar values
writer.add_scalar('Loss/Discriminator', errD.item(), global_step=iters)
writer.add_scalar('Loss/Generator', errG.item(), global_step=iters)
writer.add_scalar('Performance/D(x)', D_x, global_step=iters)
writer.add_scalar('Performance/D(G(z1))', D_G_z1, global_step=iters)
writer.add_scalar('Performance/D(G(z2))', D_G_z2, global_step=iters)

# Log images generated by the GAN
if iters % 500 == 0 or ((epoch == num_epochs-1) and (i == len(train_dataloader) - 1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
    img_grid = vutils.make_grid(fake, padding=2, normalize=True)
    writer.add_image('Generated Images', img_grid, global_step=iters)

iters += 1

```



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Starting Training Loop...

Epoch	Iteration	Loss_D	Loss_G	D(x)	D(G(z))
0	50	1.8856	1.3766	0.2853	0.37
0	50	0.0734	6.1430	0.9705	0.04
0	50	0.0200	7.0660	0.9933	0.02
0	50	0.2827	8.9739	0.8583	0.02
0	50	0.0866	5.5112	0.9695	0.02
0	50	0.1766	4.3192	0.8790	0.02
0	50	0.2166	3.4513	0.8553	0.02
0	50	0.1073	3.7556	0.9503	0.02
0	50	0.1605	3.8256	0.9367	0.02
0	50	0.5088	1.0462	0.6536	0.02
0	50	0.1870	3.3786	0.9149	0.02
0	50	0.0776	3.6571	0.9545	0.02
0	50	0.3016	2.5315	0.8422	0.02

```

[0/50] [650/1138] Loss_D: 0.1148 Loss_G: 3.2126 D(x): 0.9519 D(G(z)): 0.1148
[0/50] [700/1138] Loss_D: 0.2113 Loss_G: 3.0272 D(x): 0.9257 D(G(z)): 0.1148
[0/50] [750/1138] Loss_D: 0.2317 Loss_G: 2.3149 D(x): 0.8929 D(G(z)): 0.1148
[0/50] [800/1138] Loss_D: 0.4373 Loss_G: 3.2826 D(x): 0.9385 D(G(z)): 0.1148
[0/50] [850/1138] Loss_D: 0.1860 Loss_G: 2.9189 D(x): 0.9311 D(G(z)): 0.1148
[0/50] [900/1138] Loss_D: 0.1644 Loss_G: 3.1573 D(x): 0.9354 D(G(z)): 0.1148
[0/50] [950/1138] Loss_D: 0.2125 Loss_G: 2.7380 D(x): 0.8871 D(G(z)): 0.1148
[0/50] [1000/1138] Loss_D: 1.2272 Loss_G: 0.0016 D(x): 0.3684 D(G(z)): 0.1148
[0/50] [1050/1138] Loss_D: 0.2499 Loss_G: 2.9425 D(x): 0.8987 D(G(z)): 0.1148
[0/50] [1100/1138] Loss_D: 0.3494 Loss_G: 1.9796 D(x): 0.7954 D(G(z)): 0.1148
[1/50] [0/1138] Loss_D: 0.1809 Loss_G: 3.2387 D(x): 0.9552 D(G(z)): 0.1148
[1/50] [50/1138] Loss_D: 0.4720 Loss_G: 1.7447 D(x): 0.7897 D(G(z)): 0.1148
[1/50] [100/1138] Loss_D: 0.1476 Loss_G: 3.2569 D(x): 0.9431 D(G(z)): 0.1148
[1/50] [150/1138] Loss_D: 0.6997 Loss_G: 7.5305 D(x): 0.9899 D(G(z)): 0.1148
[1/50] [200/1138] Loss_D: 0.3494 Loss_G: 3.0406 D(x): 0.9276 D(G(z)): 0.1148
[1/50] [250/1138] Loss_D: 0.5386 Loss_G: 2.3750 D(x): 0.7172 D(G(z)): 0.1148
[1/50] [300/1138] Loss_D: 0.2177 Loss_G: 2.4332 D(x): 0.9372 D(G(z)): 0.1148
[1/50] [350/1138] Loss_D: 1.4094 Loss_G: 0.0092 D(x): 0.3663 D(G(z)): 0.1148
[1/50] [400/1138] Loss_D: 0.3489 Loss_G: 2.1725 D(x): 0.8232 D(G(z)): 0.1148
[1/50] [450/1138] Loss_D: 0.2670 Loss_G: 1.8421 D(x): 0.8479 D(G(z)): 0.1148
[1/50] [500/1138] Loss_D: 0.1895 Loss_G: 2.4232 D(x): 0.9109 D(G(z)): 0.1148
[1/50] [550/1138] Loss_D: 0.1631 Loss_G: 3.4621 D(x): 0.8747 D(G(z)): 0.1148
[1/50] [600/1138] Loss_D: 0.5265 Loss_G: 2.8027 D(x): 0.9022 D(G(z)): 0.1148
[1/50] [650/1138] Loss_D: 0.3135 Loss_G: 2.4694 D(x): 0.8936 D(G(z)): 0.1148
[1/50] [700/1138] Loss_D: 0.1190 Loss_G: 2.9570 D(x): 0.9550 D(G(z)): 0.1148
[1/50] [750/1138] Loss_D: 1.1398 Loss_G: 1.0827 D(x): 0.5845 D(G(z)): 0.1148
[1/50] [800/1138] Loss_D: 0.1905 Loss_G: 3.2324 D(x): 0.8867 D(G(z)): 0.1148
[1/50] [850/1138] Loss_D: 0.1749 Loss_G: 4.0558 D(x): 0.9545 D(G(z)): 0.1148
[1/50] [900/1138] Loss_D: 0.6403 Loss_G: 1.1777 D(x): 0.5743 D(G(z)): 0.1148
[1/50] [950/1138] Loss_D: 1.0962 Loss_G: 1.2221 D(x): 0.7071 D(G(z)): 0.1148
[1/50] [1000/1138] Loss_D: 0.4899 Loss_G: 3.9073 D(x): 0.9624 D(G(z)): 0.1148
[1/50] [1050/1138] Loss_D: 1.1566 Loss_G: 0.7766 D(x): 0.4426 D(G(z)): 0.1148
[1/50] [1100/1138] Loss_D: 0.2599 Loss_G: 2.0171 D(x): 0.8386 D(G(z)): 0.1148
[2/50] [0/1138] Loss_D: 0.1813 Loss_G: 3.4713 D(x): 0.8491 D(G(z)): 0.0000
[2/50] [50/1138] Loss_D: 0.4838 Loss_G: 2.4311 D(x): 0.8045 D(G(z)): 0.1148
[2/50] [100/1138] Loss_D: 0.2966 Loss_G: 2.0177 D(x): 0.8254 D(G(z)): 0.1148
[2/50] [150/1138] Loss_D: 0.1374 Loss_G: 2.8968 D(x): 0.8971 D(G(z)): 0.1148
[2/50] [200/1138] Loss_D: 0.2701 Loss_G: 3.1398 D(x): 0.8833 D(G(z)): 0.1148
[2/50] [250/1138] Loss_D: 0.1641 Loss_G: 3.8185 D(x): 0.9850 D(G(z)): 0.1148
[2/50] [300/1138] Loss_D: 1.0090 Loss_G: 0.6511 D(x): 0.4305 D(G(z)): 0.1148
[2/50] [350/1138] Loss_D: 0.2025 Loss_G: 2.0010 D(x): 0.8005 D(G(z)): 0.1148

```

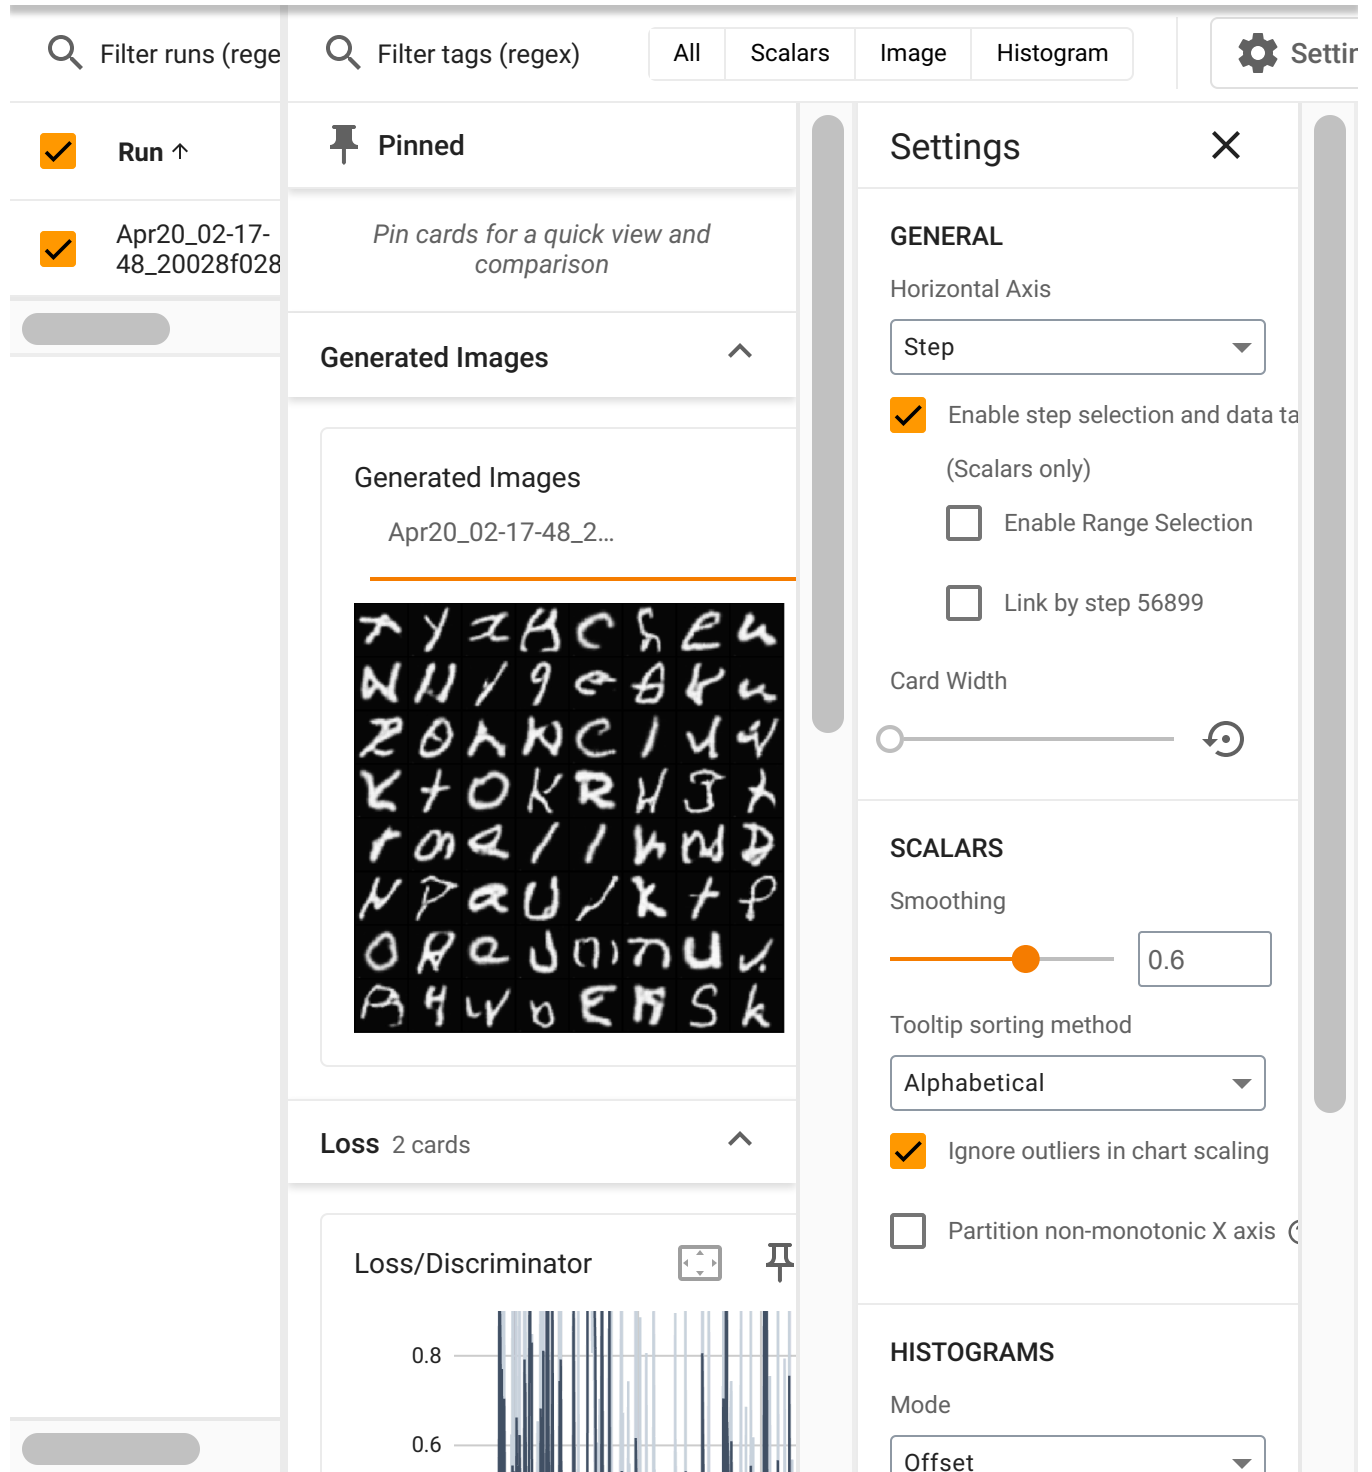
```
%tensorboard --logdir runs
```



TensorBoard


TIME SERIES

SCALARS INACTIVE



```
plt.figure(figsize=(10,5))
plt.title("Generator and Discriminator Loss During Training")
plt.plot(G_losses,label="G")
plt.plot(D_losses,label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
```

```
plt.legend()
plt.show()
```



```
fig = plt.figure(figsize=(8,8))
plt.axis("off")
ims = [[plt.imshow(np.transpose(i,(1,2,0))), animated=True] for i in img_list]
ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000, blit=True)

HTML(ani.to_jshtml())
```

⚡ WARNING:matplotlib.animation:Animation size has reached 21053899 bytes, exceeding

